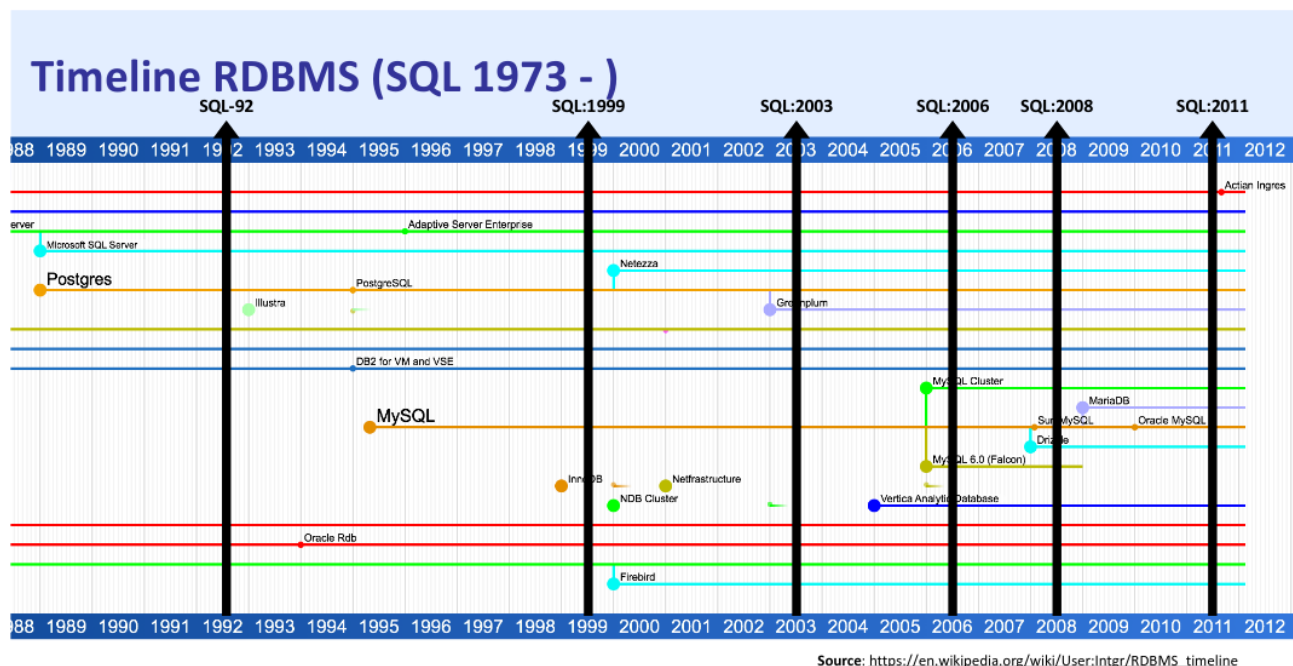


# STRUCTURED QUERY LANGUAGE

This was originally named SEQUEL (Structured English QUery Language)

- Designed and implemented by IBM Research as an interface of experimental relational database SYSTEM R
  - High-level “non-navigational” language, ad-hoc queries
  - Support rapidly changing database environment
- Used for data definition as well as queries and updates (both DDL and DML)
- SQL can also specify authorization and security, define integrity constraints, define views, specify transaction controls.



Important take aways:

- Core SQL concepts and Keywords most common in SQL compliant versions (backward compat.)
- Concept that SQL will have variation, and not to get stuck in a vendor-specific version.
- Learn how to adapt to a particular version.

## RELATIONAL MODEL TERMS AND SQL TERMS

| <b>Formal Terms<br/>(Relational Model)</b> | <b>Informal Terms<br/>(SQL)</b> |
|--|---------------------------------|
| Relation                                   | Table                           |
| Attribute                                  | Column Header                   |
| Domain                                     | All possible Column Values      |
| Tuple                                      | Row                             |
| <b>Schema</b> of a Relation                | Table Definition                |
| <b>State</b> of the Relation               | Populated Table                 |

## DATA DEFINITION, CONSTRAINTS AND SCHEMA CHANGES

Data Definition Language (DDL) Commands:

- CREATE - Create a description of the relations
- DROP - Delete the descriptions
- ALTER - Update the descriptions

### CREATE SCHEMA

- Specifies a new database schema by giving it a name.

```
CREATE SCHEMA COMPANY AUTHORIZATION JSMITH;
```

- Selects a schema to be defined:

```
USE COMPANY;
```

- Multiple schemas exist within a database, although some RDBMSs have schema and a database as a synchronous concept.
- Comments:

--line comments are two dashes

So to create the table DEPARTMENT:

```
CREATE TABLE DEPARTMENT (
    DNAME          VARCHAR(10)      NOT NULL,
    DNUMBER        INTEGER          NOT NULL,
    MGRSSN         CHAR(9),
    MGRSTARTDATE   DATE
);
```

- Specifies a new base relation by giving it a name, and specifying each of its attributes.
- In SQL, attributes are ordered based on the order they are specified.
- Attributes can have initial constraints defined, as in NOT NULL

## ATTRIBUTE NAMES AND TABLE NAMES RESERVED WORDS IN SQL

- Vendor specific... and can be a long list...
- How to deal with reserved words? Use back-ticks if you *really* want to use reserved words.

``GROUP``      `INTEGER`      `NOT NULL`

Table 9.2 Keywords and Reserved Words in MySQL 5.5

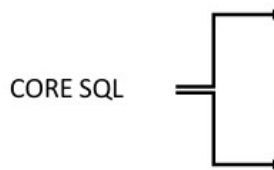
|                 |                |            |
|-----------------|----------------|------------|
| ACCESSIBLE (R)  | ACTION         | ADD (R)    |
| AFTER           | AGAINST        | AGGREGATE  |
| ALGORITHM       | ALL (R)        | ALTER (R)  |
| ANALYZE (R)     | AND (R)        | ANY        |
| AS (R)          | ASC (R)        | ASCII      |
| ASENSITIVE (R)  | AT             | AUTHORS    |
| AUTOEXTEND_SIZE | AUTO_INCREMENT | AVG        |
| AVG_ROW_LENGTH  | BACKUP         | BEFORE (R) |
| BEGIN           | BETWEEN (R)    | BIGINT (R) |
| BINARY (R)      | BINLOG         | BIT        |
| BLOB (R)        | BLOCK          | BOOL       |
| BOOLEAN         | BOTH (R)       | BTREE      |
| BY (R)          | BYTE           | CACHE      |
| CALL (R)        | CASCADE (R)    | CASCADE    |
| CASE (R)        | CATALOG_NAME   | CHAIN      |
| CHANGE (R)      | CHANGED        | CHAR (R)   |

## ATTRIBUTE DATA TYPES

### INTEGER

- INT or INTEGER are signed implicitly
- UNSIGNED INTEGER are unsigned
- INT(n) denotes the number of digits (i.e. INT(2) 0-99)
- Size depends on implementation

Example: MySQL Implementation of Integer Types



| Type      | Storage<br>(Bytes) | Minimum Value<br>(Signed/Unsigned) | Maximum Value<br>(Signed/Unsigned) |
|-----------|--------------------|------------------------------------|------------------------------------|
| TINYINT   | 1                  | -128                               | 127                                |
|           |                    | 0                                  | 255                                |
| SMALLINT  | 2                  | -32768                             | 32767                              |
|           |                    | 0                                  | 65535                              |
| MEDIUMINT | 3                  | -8388608                           | 8388607                            |
|           |                    | 0                                  | 16777215                           |
| INT       | 4                  | -2147483648                        | 2147483647                         |
|           |                    | 0                                  | 4294967295                         |
| BIGINT    | 8                  | -9223372036854775808               | 9223372036854775807                |
|           |                    | 0                                  | 18446744073709551615               |

### REAL NUMBERS

Approximate Value:

- FLOAT, REAL, DOUBLE
- Can specify digit precision (DB does *rounding*)
- FLOAT(n)
  - n – precision, as in number of bits used to store the mantissa of the float number in scientific notation
- Example: FLOAT(24) holds a single precision floating point number

Exact Value (Fixed-Point Type):

- DECIMAL(i, j) for exact formatted numbers
  - i – precision, total number of digits to store number
  - j – scale, after decimal
- Example: DECIMAL(7, 4) will look like -999.9999 when displayed

## STRING

- CHAR(n), CHARACTER(n) – Fixed length, right padded with spaces
- VARCHAR(n), CHAR VARYING(n), CHARACTER VARYING(n) – Varying length
- CLOB / TEXT – Character Large Object

| Value      | CHAR (4) | Storage Required | VARCHAR (4) | Storage Required |
|------------|----------|------------------|-------------|------------------|
| ' '        | ' '      | 4 bytes          | ' '         | 1 byte           |
| 'ab'       | 'ab '    | 4 bytes          | 'ab'        | 3 bytes          |
| 'abcd'     | 'abcd'   | 4 bytes          | 'abcd'      | 5 bytes          |
| 'abcdefgh' | 'abcd'   | 4 bytes          | 'abcd'      | 5 bytes          |

## BINARY DATA

- BIT(n) – Fixed length
- BIT VARYING(n) – Varying length
- BLOB – Binary Large. Megabyte, Gigabyte.

Typically, people do not store files (images etc.) or documents in a RDBS. They are stored on a File System with reference.

PRO TIP: Do *hash* your binary data and store that in the DB with the metadata (Use “md5 hash” of your binary data).

## BOOLEAN

- BOOLEAN – TRUE or FALSE
- Can implement with BIT(1) – 1 bit
- Implementations vary, examples:
  - MySQL uses a BIT(1), so 0 or 1 as False/True
  - PostgreSQL stores these as any of these literals →

|        |   |         |
|--------|---|---------|
| TRUE   | : | FALSE   |
| 't'    | : | 'f'     |
| 'true' | : | 'false' |
| 'y'    | : | 'n'     |
| 'yes'  | : | 'no'    |
| 'on'   | : | 'off'   |
| '1'    | : | '0'     |

*Any data type can also be NULL...*

## NULL

- Attributes in SQL can have the value of NULL (unless of course they are specified with NOT NULL constraints)
- NULL means Unknown data, and does not mean False.
- With Boolean values and Conditionals, there is a Three Value Logic System in SQL

| $p$     | $q$     | $p \text{ OR } q$ | $p \text{ AND } q$ | $p = q$ |
|---------|---------|-------------------|--------------------|---------|
| True    | True    | True              | True               | True    |
| True    | False   | True              | False              | False   |
| True    | Unknown | True              | Unknown            | Unknown |
| False   | True    | True              | False              | False   |
| False   | False   | False             | False              | True    |
| False   | Unknown | Unknown           | False              | Unknown |
| Unknown | True    | True              | Unknown            | Unknown |
| Unknown | False   | Unknown           | False              | Unknown |
| Unknown | Unknown | Unknown           | Unknown            | Unknown |

| $p$     | NOT $p$ |
|---------|---------|
| True    | False   |
| False   | True    |
| Unknown | Unknown |

## DATE AND TIME

- DATE – Made up of year-month-day (“yyyy-mm-dd”)
- TIME – Made up of hour:minute:second (“hh:mm:ss”)
- TIME(i) – TIME plus  $i$  additional digits for fractions of second
  - TIME(3) → precision for milliseconds
- DATETIME/TIMESTAMP – both DATE and TIME components

## INTERVAL

INTERVAL – relative time value as opposed to absolute

- Can be DAY/TIME intervals or YEAR/MONTH intervals
- Can be positive or negative when added to or subtracted from an absolute value, the result is an absolute value

Examples of Interval Values that can be stored: INTERVAL 1 DAY, INTERVAL 3 MONTH, INTERVAL 2 HOUR

## COMPLEX TYPES

- Special types:
  - CURRENCY, MONEY
- Spatial Types (GIS)
  - GEOMETRY type can store geometries
    - POINT(0 0)
    - LINESTRING(0 0, 1 1, 1 2)
    - POLYGON((0 0, 4 0, 4 4, 0 4, 0 0), (1 1, 2 1, 2 2, 1 2, 1 1))
- Enumerated Types: ENUM("One", "Two", "Three")
- Collection Types: SET, VALUE\_MAP

## CREATING DOMAINS

- CREATE DOMAIN - Allows you to specify your own data type to use in schema  

```
CREATE DOMAIN SSN AS CHAR(9);
```
- Need to CREATE DOMAIN before utilizing in CREATE TABLE

## SPECIFYING INTEGRITY CONSTRAINTS

- NOT NULL - for enforcing that attributes cannot take NULL values  

```
DNUMBER INTEGER NOT NULL;
```
- DEFAULT <value>  

```
MGRSSN CHAR(9) DEFAULT '123456789';
```
- AUTO\_INCREMENT - For INTEGER types, helpful with IDs (usually starts with 1, but check implementation...)  

```
DNUMBER INTEGER NOT NULL AUTO_INCREMENT
```

## CHECK

- CHECK clause - requires a valid conditional expression  

```
DNUMBER INT CHECK (DNUMBER > 0 AND DNUMBER < 21);
```
- Use with CREATE DOMAIN (note the use of VALUE to reference the attribute name):  

```
CREATE DOMAIN D_NUM AS INTEGER CHECK (VALUE > 0 AND VALUE < 21);
```
- Check is unable to compare against other attributes/relations - need to use an ASSERTION

## KEY AND REFERENTIAL INTEGRITY CONSTRAINTS

- PRIMARY KEY clause

DNUMBER INT PRIMARY KEY;

- UNIQUE clause – for secondary/alternate keys

DNAME CHAR(9) UNIQUE;

- FOREIGN KEY clause – for referential integrity

FOREIGN KEY (MGRSSN) REFERENCES EMPLOYEE(SSN);

## INTEGRITY CONSTRAINTS

```
CREATE TABLE DEPARTMENT (
    DNAME          VARCHAR(10)      NOT NULL,
    DNUMBER        INTEGER          NOT NULL,
    MGRSSN         CHAR(9) ,
    MGRSTARTDATE   DATE ,
    PRIMARY KEY (DNUMBER) ,
    UNIQUE (DNAME) ,
    FOREIGN KEY (MGRSSN) REFERENCES EMPLOYEE (SSN)
);
```

- CREATE TABLE can specify primary key attributes, secondary keys, and referential integrity constraints (foreign keys) after the attributes.
- Key attributes specified via PRIMARY KEY and UNIQUE

## COMPOSITE PRIMARY AND FOREIGN KEYS

Composite PRIMARY or FOREIGN KEYS are specified after the attributes

```
CREATE TABLE DEPARTMENT (
    DNAME          VARCHAR(10)      UNIQUE NOT NULL,
    DNUMBER        INTEGER          NOT NULL,
    MGR_FNAME      CHAR(9) ,
    MGR_LNAME      CHAR(9) ,
    MGRSTARTDATE   DATE ,
    PRIMARY KEY (DNUMBER, DNAME) ,
    FOREIGN KEY (MGR_FNAME, MGR_LNAME) REFERENCES EMPLOYEE (FNAME, LNAME)
);
```



## REFERENTIAL INTEGRITY OPTIONS

- A referential integrity constraint may be violated when tuples in the referenced tuple are updated/deleted.
- Default: reject the operation that violates constraint
- Or: Specify a referential triggered action, options:

**Events**  
(what occurs on  
Referenced Tuple with PK)

ON DELETE

ON UPDATE

**Triggered Action**  
(to happen on Referencing Tuple with FK)

RESTRICT

CASCADE

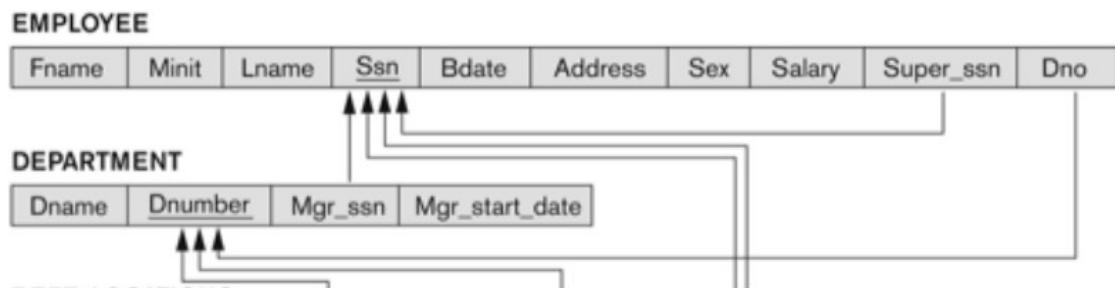
SET NULL

SET DEFAULT

Example:

ON DELETE SET NULL  
ON UPDATE CASCADE

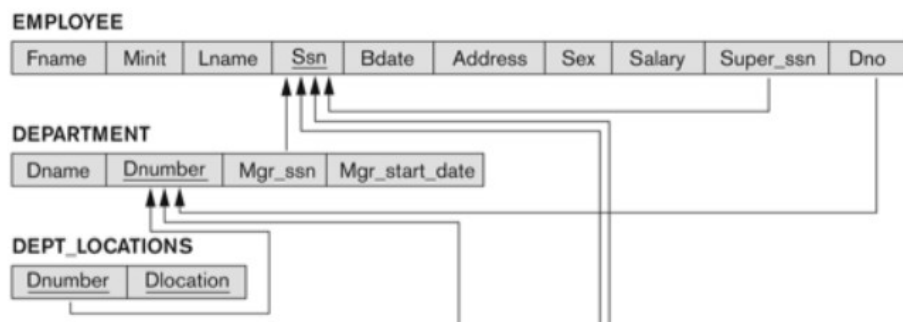
```
CREATE TABLE DEPARTMENT (
...
    FOREIGN KEY (MGRSSN) REFERENCES EMPLOYEE (SSN)
    ON DELETE SET DEFAULT ON UPDATE CASCADE
...
);
```



```

CREATE TABLE EMPLOYEE (
...
    FOREIGN KEY(DNO) REFERENCES DEPARTMENT(DNUMBER) ON DELETE SET
    DEFAULT ON UPDATE CASCADE,
    FOREIGN KEY(SUPERSSN) REFERENCES EMPLOYEE(SSN) ON DELETE SET
    NULL ON UPDATE CASCADE,
...
);

```



## BASE RELATIONS VS. VIRTUAL RELATIONS

- Base tables (or base relations): Relations and tuples store as a file by DBMS
  - Created through the CREATE TABLE statement
- Base tables are distinguished from virtual relations which may not correspond to an actual physical file
  - Created through the CREATE VIEW statement
- Remember: while attributes in CREATE TABLE are ordered, tuples (rows) are not considered ordered.



## DROP TABLE

- Used to remove a relation (base table) and its definition
- Relation is unable to be used in queries, updates, or any commands since its description no longer exists
- Example

```
DROP TABLE DEPENDENT;
```

- Table dropped if not referenced in any constraints

```
DROP TABLE DEPENDENT RESTRICT;
```

- All constraints that reference table are dropped along with table.

```
DROP TABLE DEPENDENT CASCADE;
```

## DROP SCHEMA

- Used to remove the entire schema
- Similar distinction with RESTRICT vs. CASCADE as with DROP TABLE
- Dropped only if no elements in schema

```
DROP SCHEMA COMPANY RESTRICT;
```

- ALL tables, views, and constraints dropped (!!!)

```
DROP SCHEMA DEPENDENT CASCADE;
```

## ALTER TABLE

### ADD

Used to add an attribute to one of the base relations

- New attribute will have NULLS in all existing tuples of relation

```
ALTER TABLE EMPLOYEE ADD JOB VARCHAR(12);
```

- Database users will need to UPDATE a value for the new JOB attribute for the existing Employees.
- Utilize a DEFAULT value

```
ALTER TABLE EMPLOYEE ADD JOB VARCHAR(12) DEFAULT 'President';
```

- Specifying NOT NULL will require a DEFAULT value.

## ADD CONSTRAINTS

Depending on which order tables are created, circular Referential Integrity Constraints may need to be added later.

- Example, DEPARTMENT and EMPLOYEE reference each other:

```
ALTER TABLE EMPLOYEE ADD FOREIGN KEY (DNO) REFERENCES DEPARTMENT (Dnumber);
```

```
ALTER TABLE DEPARTMENT ADD FOREIGN KEY (MGRSSN) REFERENCES EMPLOYEE (Ssn);
```

## DROP

- Can remove attributes (which removes data), although some RDBS do not allow removing columns

```
ALTER TABLE EMPLOYEE DROP JOB;
```

- Need to specify the FOREIGN KEY to DROP a constraint.

```
ALTER TABLE DEPARTMENT DROP FOREIGN KEY (MGRSSN);
```

# DATA MANIPULATION LANGUAGE

---

## RETRIEVAL QUERIES IN SQL

### BAGS VS. SETS

- A bag or multi-set is like a set, but an element may appear more than once.

{A, B, C, A} is a bag.

{A, B, C} is also a bag that is also a set

- Bags also resemble lists, but order is irrelevant in a bag.

{A, B, A} = {B, A, A} as bags

However, [A, B, A]  $\neq$  [B, A, A] as lists

- SQL can enforce sets with Key Constraints and DISTINCT
- Ordered relations (lists) with ORDER BY

### SELECT FROM WHERE

- Basic form of the SQL SELECT statement is called a mapping, or a SELECT-FROM-WHERE block

```
SELECT    <attribute list>
FROM      <table list>
WHERE     <condition>
```

- <attribute list> is a list of attribute names whose values are to be retrieved by the query
- <table list> is a list of the relation names required to process the query
- <condition> is a conditional (boolean) expression that identifies the tuples to be retrieved by the query

### WHERE CONDITIONS

Basic Logical Operators in the <condition> : =, <, <=, >, >=, <>, !=, AND, OR, NOT, IS NULL, IS NOT NULL

WHERE DNO = 5

WHERE (DNO = 5) AND (PNUMBER > 1)

## SIMPLE SQL QUERIES AND RELATIONAL ALGEBRA

- Basic SQL queries correspond to using the following operations of the relational algebra
  - SELECT operator
  - PROJECT operator
  - JOIN

### SIMPLE QUERY ON ONE RELATION

- Retrieve the birthdate and address of the employee whose name is 'John B. Smith'.

```
SELECT BDATE, ADDRESS
FROM EMPLOYEE
WHERE FNAME="John" AND MINIT="B" AND LNAME="Smith"
```

| EMPLOYEE | FNAME    | MINIT | LNAME   | SSN       | BDATE      | ADDRESS                  | SEX | SALARY | SUPERSSN  | DNO |
|----------|----------|-------|---------|-----------|------------|--------------------------|-----|--------|-----------|-----|
|          | John     | B     | Smith   | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M   | 30000  | 333445555 | 5   |
|          | Franklin | T     | Wong    | 333445555 | 1955-12-08 | 638 Voss, Houston, TX    | M   | 40000  | 888665555 | 5   |
|          | Alicia   | J     | Zelaya  | 999887777 | 1968-07-19 | 3321 Castle, Spring, TX  | F   | 25000  | 987654321 | 4   |
|          | Jennifer | S     | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX  | F   | 43000  | 888665555 | 4   |
|          | Ramesh   | K     | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M   | 38000  | 333445555 | 5   |
|          | Joyce    | A     | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX   | F   | 25000  | 333445555 | 5   |
|          | Ahmad    | V     | Jabbar  | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX  | M   | 25000  | 987654321 | 4   |
|          | James    | E     | Borg    | 888665555 | 1937-11-10 | 450 Stone, Houston, TX   | M   | 55000  | null      | 1   |

```
SELECT BDATE, ADDRESS
FROM EMPLOYEE
WHERE FNAME='John' AND MINIT='B'
AND LNAME='Smith'
```

*Result*

| BDATE      | ADDRESS                 |
|------------|-------------------------|
| 09/01/1965 | 731 Fondren, Houston TX |

## SIMPLE QUERY WITH MORE THAN ONE RELATION

- Retrieve the name and address of all employees who work for the 'Research' department.

```
SELECT FNAME, LNAME, ADDRESS
FROM EMPLOYEE, DEPARTMENT
WHERE DNAME = 'RESEARCH' AND DNUMBER = DNO
```

```
SELECT FNAME, LNAME, ADDRESS
FROM EMPLOYEE, DEPARTMENT
WHERE DNAME='Research' AND DNUMBER=DNO
```

*Result*

| EMPLOYEE | FNAME | MINIT   | LNAME     | SSN        | BOATE | ADDRESS                  | SEX | SALARY | SUPERSSN  | DNO |
|----------|-------|---------|-----------|------------|-------|--------------------------|-----|--------|-----------|-----|
| John     | B     | Smith   | 123456789 | 1965-01-09 |       | 731 Fondren, Houston, TX | M   | 30000  | 333445555 | 5   |
| Franklin | T     | Wong    | 333445555 | 1955-12-08 |       | 638 Voss, Houston, TX    | M   | 40000  | 888665555 | 5   |
| Alicia   | J     | Zelaya  | 999887777 | 1968-07-19 |       | 3321 Castle, Spring, TX  | F   | 25000  | 987654321 | 4   |
| Jennifer | S     | Wallace | 987654321 | 1941-06-20 |       | 291 Berry, Bellaire, TX  | F   | 43000  | 888665555 | 4   |
| Ramesh   | K     | Narayan | 666884444 | 1962-09-15 |       | 975 Fire Oak, Humble, TX | M   | 38000  | 333445555 | 5   |
| Joyce    | A     | English | 453453453 | 1972-07-31 |       | 5631 Rice, Houston, TX   | F   | 25000  | 333445555 | 5   |
| Ahmad    | V     | Jabbar  | 987987987 | 1969-03-29 |       | 980 Dallas, Houston, TX  | M   | 25000  | 987654321 | 4   |
| James    | E     | Borg    | 888665555 | 1937-11-10 |       | 450 Stone, Houston, TX   | M   | 55000  | null      | 1   |

| DEPARTMENT | DNAME          | DNUMBER | MGRSSN    | MGRSTARTDATE |
|------------|----------------|---------|-----------|--------------|
|            | Research       | 5       | 333445555 | 1988-05-22   |
|            | Administration | 4       | 987654321 | 1995-01-01   |
|            | Headquarters   | 1       | 888665555 | 1981-06-19   |

| FNAME    | LNAME   | ADDRESS                 |
|----------|---------|-------------------------|
| John     | Smith   | 731 Fondren, Houston TX |
| Franklin | Wong    | 638 Voss, Houston TX    |
| Joyce    | English | 5631 Rice, Houston TX   |
| Ramesh   | Narayan | 975 Fire Oak, Humble TX |

## THE JOIN CONDITION

What happens if the condition DNUMBER = DNO is omitted?

```
SELECT FNAME, LNAME, ADDRESS
FROM EMPLOYEE, DEPARTMENT
WHERE DNAME='Research' AND DNUMBER=DNO
```

| EMPLOYEE | FNAME | MINIT   | LNAME     | SSN        | BDATE                    | ADDRESS | SEX   | SALARY    | SUPERSSN | DNO |
|----------|-------|---------|-----------|------------|--------------------------|---------|-------|-----------|----------|-----|
| John     | B     | Smith   | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M       | 30000 | 333445555 | 5        |     |
| Franklin | T     | Wong    | 333445555 | 1955-12-08 | 638 Voss, Houston, TX    | M       | 40000 | 888665555 | 5        |     |
| Alicia   | J     | Zelaya  | 999887777 | 1968-07-19 | 3321 Castle, Spring, TX  | F       | 25000 | 987654321 | 4        |     |
| Jennifer | S     | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX  | F       | 43000 | 888665555 | 4        |     |
| Ramesh   | K     | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M       | 38000 | 333445555 | 5        |     |
| Joyce    | A     | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX   | F       | 25000 | 333445555 | 5        |     |
| Ahmad    | V     | Jabbar  | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX  | M       | 25000 | 987654321 | 4        |     |
| James    | E     | Borg    | 888665555 | 1937-11-10 | 450 Stone, Houston, TX   | M       | 55000 | null      | 1        |     |

| DEPARTMENT | DNAME          | DNUMBER | MGRSSN    | MGRSTARTDATE |
|------------|----------------|---------|-----------|--------------|
|            | Research       | 5       | 333445555 | 1988-05-22   |
|            | Administration | 4       | 987654321 | 1995-01-01   |
|            | Headquarters   | 1       | 888665555 | 1981-06-19   |

All rows are selected

| FNAME    | LNAME   | ADDRESS                 |
|----------|---------|-------------------------|
| John     | Smith   | 731 Fondren, Houston TX |
| Franklin | Wong    | 638 Voss, Houston TX    |
| Joyce    | English | 5631 Rice, Houston TX   |
| Ramesh   | Narayan | 975 Fire Oak, Humble TX |
| James    | Borg    | 450 Stone, Houston TX   |
| Jennifer | Wallace | 291 Berry, Bellaire TX  |
| Ahmad    | Jabbar  | 980 Dallas, Houston TX  |
| Alicia   | Zelaya  | 3321 Castle, Spring TX  |

## SELECT FROM MORE RELATIONS

- For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birthdate

```
SELECT PNUMBER, DNUM, LNAME, ADDRESS, BDATE
FROM PROJECT, DEPARTMENT, EMPLOYEE
WHERE DNUM=DNUMBER AND MGRSSN=SSN AND PLOCATION='STAFFORD'
```

Two Join conditions:

- The join condition DNUM=DNUMBER relates to a project to its controlling department.
- The join condition MGRSSN=SSN relates the controlling department to the employee who manages that department.

```
SELECT PNUMBER, DNUM,
       LNAME, ADDRESS, BDATE
FROM PROJECT, DEPARTMENT,
     EMPLOYEE
WHERE DNUM=DNUMBER AND
      MGRSSN=SSN AND
      PLOCATION='Stafford'
```

| EMPLOYEE | FNAME | MINIT   | LNAME     | SSN        | BDATE | ADDRESS                  | SEX | SALARY | SUPERSSN  | DNO |
|----------|-------|---------|-----------|------------|-------|--------------------------|-----|--------|-----------|-----|
| John     | B     | Smith   | 132456789 | 1965-01-09 |       | 731 Fondren, Houston, TX | M   | 30000  | 333445555 | 5   |
| Franklin | T     | Wong    | 333445555 | 1955-12-08 |       | 638 Voss, Houston, TX    | M   | 40000  | 888665555 | 5   |
| Alicia   | J     | Zelaya  | 999807777 | 1968-07-19 |       | 3321 Castle, Spring, TX  | F   | 25000  | 887654321 | 4   |
| Jennifer | S     | Wallace | 887654321 | 1941-06-20 |       | 291 Berry, Bellaire, TX  | F   | 43000  | 888665555 | 4   |
| Ramesh   | K     | Narayan | 688984444 | 1962-06-15 |       | 975 Fire Oak, Humble, TX | M   | 38000  | 333445555 | 5   |
| Joyce    | A     | English | 434353433 | 1972-07-01 |       | 5631 Rice, Houston, TX   | F   | 25000  | 333445555 | 5   |
| Ahmad    | V     | Jabbar  | 887654321 | 1969-03-29 |       | 980 Dallas, Houston, TX  | M   | 25000  | 887654321 | 4   |
| James    | E     | Borg    | 888665555 | 1957-11-10 |       | 450 Stone, Houston, TX   | M   | 55000  | null      | 1   |

| DEPT_LOCATIONS |  | DNUMBER | DLOCATION |
|----------------|--|---------|-----------|
|                |  | 1       | Houston   |
|                |  | 4       | Stafford  |
|                |  | 5       | Bellaire  |
|                |  | 5       | Superfund |
|                |  | 5       | Houston   |

| DEPARTMENT     | DNAME | DNUMBER | MGRSSN    | MGRSTARTDATE |
|----------------|-------|---------|-----------|--------------|
| Research       |       | 5       | 333445555 | 1969-05-22   |
| Administration |       | 4       | 887654321 | 1955-01-01   |
| Headquarters   |       | 1       | 888665555 | 1981-09-19   |

| PNUMBER | DNUM | LNAME   | ADDRESS                | BDATE      |
|---------|------|---------|------------------------|------------|
| 10      | 4    | Wallace | 291 Berry, Bellaire TX | 20/06/1941 |
| 30      | 4    | Wallace | 291 Berry, Bellaire TX | 20/06/1941 |



## QUALIFICATION OF RELATION NAMES

- In SQL, we can use the same name for multiple attributes as long as the attributes are in different relations
- If two or more attributes in different relations have the same name, we need to specify them by the relation name.
- We can qualify the attribute name with the relation name by prefixing the relation name to the attribute name.

E.g. Unique attribute names

```
SELECT EMPLOYEE.FNAME, EMPLOYEE.LNAME, EMPLOYEE.ADDRESS
FROM EMPLOYEE, DEPARTMENT
WHERE DEPARTMENT.DNAME='Research' AND DEPARTMENT.DNUMBER=EMPLOYEE.DNO
```

## ALIASES

- Some queries need to refer to the same relation twice
- Aliases can be given to the relation names

For each employee, retrieve the employee's name, and the name of his or her immediate supervisor.

```
SELECT E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM EMPLOYEE E, EMPLOYEE S
WHERE E.SUPERSSN=S.SSN
```

- Alternate relation names E and S are called aliases or tuple variables for the EMPLOYEE relation
- Think of E and S as two different copies of EMPLOYEE; E represents employees in role of supervisees and S represents employees in role of supervisors

| EMPLOYEE | FNAME | MINIT   | LNAME     | SSN        | BDATE                    | ADDRESS | SEX   | SALARY    | SUPERSSN | DNO |
|----------|-------|---------|-----------|------------|--------------------------|---------|-------|-----------|----------|-----|
| John     | B     | Smith   | 123456789 | 1955-01-09 | 731 Fandren, Houston, TX | M       | 30000 | 333445555 | 5        |     |
| Franklin | T     | Wong    | 333445555 | 1955-12-08 | 638 Voss, Houston, TX    | M       | 40000 | 888995555 | 5        |     |
| Alice    | J     | Zelaya  | 999887777 | 1968-07-19 | 3321 Castle, Spring, TX  | F       | 25000 | 987654321 | 4        |     |
| Jennifer | S     | Wallace | 987654321 | 1941-09-20 | 291 Berry, Dallas, TX    | F       | 43000 | 888995555 | 4        |     |
| Ramesh   | K     | Narayan | 666884444 | 1962-09-15 | 975 Pine Oak, Humble, TX | M       | 38000 | 333445555 | 5        |     |
| Joyce    | A     | English | 453453453 | 1972-07-31 | 5831 Rice, Houston, TX   | F       | 25000 | 333445555 | 5        |     |
| Ahmad    | V     | Jabbar  | 987657890 | 1969-03-29 | 980 Dallas, Houston, TX  | M       | 25000 | 987654321 | 4        |     |
| James    | E     | Borg    | 888995555 | 1937-11-10 | 450 Stone, Houston, TX   | M       | 55000 | null      | 1        |     |

```
SELECT E.FNAME AS EMPLOYEE_FNAME ,
       E.LNAME AS EMPLOYEE_LNAME ,
       S.FNAME AS SUPER_FNAME,
       S.LNAME AS SUPER_LNAME
FROM EMPLOYEE E, EMPLOYEE S
WHERE E.SUPERSSN=S.SSN
```

| EMPLOYEE_FNAME | EMPLOYEE_LNAME | SUPER_FNAME | SUPER_LNAME |
|----------------|----------------|-------------|-------------|
| John           | Smith          | Franklin    | Wong        |
| Franklin       | Wong           | James       | Borg        |
| Joyce          | English        | Franklin    | Wong        |
| Ramesh         | Narayan        | Franklin    | Wong        |
| Jennifer       | Wallace        | James       | Borg        |
| Ahmad          | Jabbar         | Jennifer    | Wallace     |
| Alicia         | Zelaya         | Jennifer    | Wallace     |

## UNSPECIFIED WHERE CLAUSE

- A missing WHERE-clause indicates no condition – all tuples of the relations in the FROM-clause are selected.

**Retrieve the SSN values for all employees,**

```
SELECT  SSN
FROM    EMPLOYEE
```

| EMPLOYEE | FNAME    | MINIT | LNAME   | SSN       | BDATE      | ADDRESS                  | SEX | SALARY | SUPERSSN  | DNO |
|----------|----------|-------|---------|-----------|------------|--------------------------|-----|--------|-----------|-----|
|          | John     | B     | Smith   | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M   | 30000  | 333445555 | 5   |
|          | Franklin | T     | Wong    | 333445555 | 1955-12-08 | 638 Voss, Houston, TX    | M   | 40000  | 888665555 | 5   |
|          | Alicia   | J     | Zelaya  | 999887777 | 1968-07-19 | 3321 Castle, Spring, TX  | F   | 25000  | 987654321 | 4   |
|          | Jennifer | S     | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX  | F   | 43000  | 888665555 | 4   |
|          | Ramesh   | K     | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M   | 38000  | 333445555 | 5   |
|          | Joyce    | A     | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX   | F   | 25000  | 333445555 | 5   |
|          | Ahmad    | V     | Jabbar  | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX  | M   | 25000  | 987654321 | 4   |
|          | James    | E     | Borg    | 888665555 | 1937-11-10 | 450 Stone, Houston, TX   | M   | 55000  | null      | 1   |

*Result*

| SSN       |
|-----------|
| 123456789 |
| 333445555 |
| 999887777 |
| 987654321 |
| 666884444 |
| 453453453 |
| 987987987 |
| 888665555 |

## CARTESIAN PRODUCT

- If more than one relation is specified in the FROM-clause and there is no join condition, then the CARTESIAN PRODUCT of tuples is selected.

```
SELECT  SSN, DNAME
FROM    EMPLOYEE, DEPARTMENT
```

*Result*

| EMPLOYEE | FNAME | MINT    | LNAME     | SSN        | BDATE                    | ADDRESS | SEX   | SALARY    | SUPERSSN | DNO |
|----------|-------|---------|-----------|------------|--------------------------|---------|-------|-----------|----------|-----|
| John     | B     | Smith   | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M       | 30000 | 333445555 | 5        |     |
| Franklin | T     | Wong    | 333445555 | 1955-12-08 | 638 Voss, Houston, TX    | M       | 40000 | 888665555 | 5        |     |
| Alicia   | J     | Zelaya  | 999887777 | 1968-07-19 | 3321 Castle, Spring, TX  | F       | 25000 | 987654321 | 4        |     |
| Jennifer | S     | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX  | F       | 43000 | 888665555 | 4        |     |
| Ramesh   | K     | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M       | 38000 | 333445555 | 5        |     |
| Joyce    | A     | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX   | F       | 25000 | 333445555 | 5        |     |
| Ahmad    | V     | Jabbar  | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX  | M       | 25000 | 987654321 | 4        |     |
| James    | E     | Borg    | 888665555 | 1937-11-10 | 450 Stone, Houston, TX   | M       | 55000 | 0         | 1        |     |

| DEPARTMENT     | DNAME | DNUMBER   | MGRSSN     | MGRSTARTDATE |
|----------------|-------|-----------|------------|--------------|
| Research       | 5     | 333445555 | 1968-05-22 |              |
| Administration | 4     | 987654321 | 1965-01-01 |              |
| Headquarters   | 1     | 888665555 | 1961-06-19 |              |

*Be careful: Easy to get LARGE relations as a result!*

| SSN       | DNAME          |
|-----------|----------------|
| 888665555 | Administration |
| 888665555 | Headquarters   |
| 888665555 | Research       |
| 987654321 | Administration |
| 987654321 | Headquarters   |
| 987654321 | Research       |
| 987987987 | Administration |
| 987987987 | Headquarters   |
| 987987987 | Research       |
| 999887777 | Administration |

## USE OF \*

To retrieve all values in tuples, a \* is used.

```
SELECT *
FROM EMPLOYEE
WHERE DNO=5
```

| FNAME    | MINIT | LNAME   | SSN       | BDATE      | ADDRESS                 | SEX | SALARY | SUPERSSN  | DNO |
|----------|-------|---------|-----------|------------|-------------------------|-----|--------|-----------|-----|
| John     | B     | Smith   | 123456789 | 09/01/1965 | 731 Fondren, Houston TX | M   | 30000  | 333445555 | 5   |
| Franklin | T     | Wong    | 333445555 | 08/12/1965 | 638 Voss, Houston TX    | M   | 40000  | 888665555 | 5   |
| Joyce    | A     | English | 453453453 | 31/07/1972 | 5631 Rice, Houston TX   | F   | 25000  | 333445555 | 5   |
| Ramesh   | K     | Narayan | 666884444 | 15/09/1962 | 975 Fire Oak, Humble TX | M   | 38000  | 333445555 | 5   |

## USE OF DISTINCT

- SQL does not treat a relation as a set; duplicate values can appear
- To eliminate duplicate tuples in a query result, the keyword DISTINCT is used

**SELECT ALL**  
SALARY  
FROM  
EMPLOYEE

| SALARY |
|--------|
| 38000  |
| 43000  |
| 25000  |
| 25000  |

**SELECT DISTINCT**  
SALARY  
FROM  
EMPLOYEE

| SALARY |
|--------|
| 38000  |
| 43000  |
| 25000  |

- Using DISTINCT with more than one attribute, creates a SET of the resulting tuples.

```
SELECT DISTINCT SEX, SALARY
FROM      EMPLOYEE
```

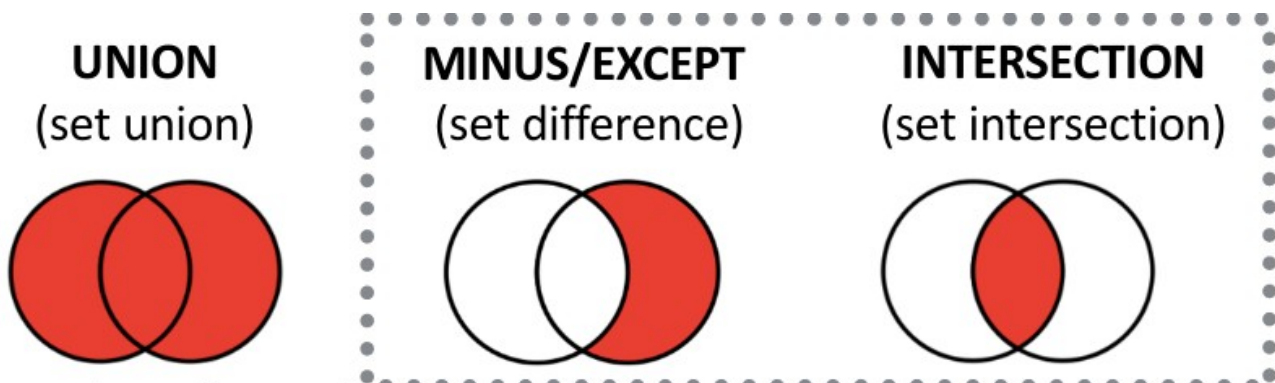
*Result: one duplicate row is removed*

| SEX          | SALARY           |
|--------------|------------------|
| M            | 30000            |
| M            | 40000            |
| F            | 25000            |
| F            | 43000            |
| M            | 38000            |
| <del>F</del> | <del>25000</del> |
| M            | 25000            |
| M            | 55000            |



## SET OPERATIONS

SQL has directly incorporated some set operations



- Resulting relations of these set operations are sets of tuples – duplicate tuples are eliminated from the result
- Set operations apply only to union compatible relations:
  - Two relations must have the same number of attributes
  - Each corresponding pair of attributes has the same domain

## UNION

- Make a list of all project numbers for projects that involve an employee whose last name is 'Smith' as a worker or as a manager of the department that controls the project.

```
(SELECT PNUMBER
FROM PROJECT, DEPARTMENT, EMPLOYEE
WHERE DNUM = DNUMBER AND MGRSSN=SSN AND LNAME = 'SMITH')
UNION
(SELECT PNUMBER
FROM PROJECT, WORKS_ON, EMPLOYEE
WHERE PNUMBER = PNO AND ESSN = SSN AND LNAME = 'SMITH')
```

## SUBSTRING COMPARISON/PATTERN MATCHING

- The LIKE comparison operator is used to compare partial strings
- Two reserved characters are used:
  - '%' (or '\*' in some implementations) replaces an arbitrary number of characters
  - '\_' replaces a single arbitrary character
- Usually use an escape-character '\' to specify these reserve characters in your search string
  - LIKE '%15\%%'

Retrieve all employees whose address is in Houston, Texas. Here, the value of the ADDRESS attribute must contain the substring 'Houston, TX' in it.

```
SELECT FNAME, LNAME
FROM EMPLOYEE
WHERE ADDRESS LIKE '%Houston, TX%'
```

| EMPLOYEE | FNAME    | MINIT | LNAME   | SSN       | BDATE      | ADDRESS                  | SEX | SALARY | SUPERSSN  | DNO |
|----------|----------|-------|---------|-----------|------------|--------------------------|-----|--------|-----------|-----|
|          | John     | B     | Smith   | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M   | 30000  | 333445555 | 5   |
|          | Franklin | T     | Wong    | 333445555 | 1955-12-08 | 638 Volts, Houston, TX   | M   | 40000  | 888665555 | 5   |
|          | Alicia   | J     | Zelaya  | 999887777 | 1968-07-19 | 3321 Castle, Spring, TX  | F   | 25000  | 987654321 | 4   |
|          | Jennifer | S     | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX  | F   | 43000  | 888665555 | 4   |
|          | Ramesh   | K     | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M   | 38000  | 333445555 | 5   |
|          | Joyce    | A     | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX   | F   | 25000  | 333445555 | 5   |
|          | Ahmad    | V     | Jabbar  | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX  | M   | 25000  | 987654321 | 4   |
|          | James    | E     | Borg    | 888665555 | 1937-11-10 | 450 Stone, Houston, TX   | M   | 55000  | null      | 1   |

Retrieve all employees who were born during the 1950s

```
SELECT FNAME, LNAME
FROM EMPLOYEE
WHERE BDATE LIKE '___5____'
```

- LIKE operator different from the formal relational model which considers each attribute value as atomic and indivisible.

```
SELECT FNAME, LNAME
FROM EMPLOYEE
WHERE BDATE LIKE '___5_____'
```

| EMPLOYEE | FNAME    | MINIT | LNAME   | SSN       | BDATE      | ADDRESS                  | SEX | SALARY | SUPERSSN  | DNO |
|----------|----------|-------|---------|-----------|------------|--------------------------|-----|--------|-----------|-----|
|          | John     | B     | Smith   | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M   | 30000  | 333445555 | 5   |
|          | Franklin | T     | Wong    | 333445555 | 1955-12-08 | 638 Voss, Houston, TX    | M   | 40000  | 888665555 | 5   |
|          | Alicia   | J     | Zelaya  | 999887777 | 1968-07-19 | 3321 Castle, Spring, TX  | F   | 25000  | 987654321 | 4   |
|          | Jennifer | S     | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX  | F   | 43000  | 888665555 | 4   |
|          | Ramesh   | K     | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M   | 38000  | 333445555 | 5   |
|          | Joyce    | A     | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX   | F   | 25000  | 333445555 | 5   |
|          | Ahmad    | V     | Jabbar  | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX  | M   | 25000  | 987654321 | 4   |
|          | James    | E     | Borg    | 888665555 | 1937-11-10 | 450 Stone, Houston, TX   | M   | 55000  | null      | 1   |

## ARITHMETIC OPERATIONS

- The standard arithmetic operators +, -, \*, and / can be applied to numeric values in an SQL query result
- Show the effect of giving all employees who work on the 'ProductX' project a 10% raise.

```
SELECT FNAME, LNAME, 1.1*SALARY
FROM EMPLOYEE, WORKS_ON, PROJECT
WHERE SSN=ESSN AND PNO=PNUMBER AND PNAME='ProductX'
```

## ORDER BY

The ORDER BY clause is used to sort the tuples in a query result based on the values of some attribute(s)

Retrieve a list of employees and the projects each works in, ordered by the employee's department, and within each department ordered alphabetically by employee last name.

```
SELECT  DNAME, LNAME, FNAME, PNAME
FROM    DEPARTMENT, EMPLOYEE,
        WORKS_ON, PROJECT
WHERE   DNUMBER=DNO AND SSN=ESSN
        AND PNO=PNUMBER
ORDER BY DNAME, LNAME
```

- The default order is in ascending order of values.
- We can specify the keyword DESC if we want a descending order; the keyword ASC can be used to explicitly specify ascending order, even though it is the default

```
SELECT  DNAME, LNAME, FNAME, PNAME
FROM    DEPARTMENT, EMPLOYEE,
        WORKS_ON, PROJECT
WHERE   DNUMBER=DNO AND SSN=ESSN
        AND PNO=PNUMBER
ORDER BY DNAME DESC, LNAME ASC
```